

Secure communication channel configurations
Validation with a namespace-based infrastructure

Introduction

This guide shows how to emulate a **two-site VPN** by combining **WSL2**, Linux **network namespaces**, and some state-of-the-art implementations for channel protection: **XFRM**, **StrongSwan**, and **OpenVPN**. The purpose is to *validate secure-channel configurations* by observing what happens “on the wire” and applying real-world WAN impairments and characteristics to bring the infrastructure closer to a real **cloud** scenario.

- **Namespaces.** Four network namespaces are implemented:
 - *hA*, *hB*: two ordinary hosts *10.0.1.2* (Host-A) and *10.0.2.2* (Host-B) in two subnets *10.0.1.0/24* (LAN-A) and *10.0.2.0/24* (LAN-B);
 - *gwA*, *gwB*: two VPN gateways, each with a LAN-side interface *10.0.1.1* (GW-A) and *10.0.2.1* (GW-B) and a WAN-side interface *192.0.2.1* (GW-A) and *192.0.2.2* (GW-B)
- **Virtual cabling.** Three veth pairs create the physical layout entirely in RAM:

Host-A (hA) ↔ GW-A (gwA) GW-A (gwA) ↔ GW-B (gwB) GW-B (gwB) ↔ Host-B (hB)




Figure 1: High-level network topology

0 Prerequisites (one-time)

WSL 2 is required, and it is suggested to have the most recent stable version of the kernel.

```
# Windows shell: check WSL/kernel
wsl --version          # must print "version 2.x.x.x"
```

Check and eventually install every tool used later in the guide: *iproute2* for namespaces, *tcpdump* for packet checks, *iperf3* for throughput tests, and VPN stacks (strongSwan, OpenVPN).

```
# WSL shell
sudo apt update
sudo apt install -y iproute2 tcpdump iperf3 strongswan strongswan-swanctl openvpn
```

Setting *systemd=true* converts the WSL distro from a minimal container-style environment into one that behaves like a normal Linux host with *PID=1*, which is what daemons such as charon-systemd (and many others) expect.

```
sudo tee /etc/wsl.conf <<'EOF'
[boot]
systemd=true
EOF
# Windows shell
wsl --shutdown # reboot the VM
# Back in WSL shell
ps -p 1 -o comm= # should print "systemd"
```

1 Namespaces setup

1.0 Target topology

Four namespaces simulate two LAN hosts and two IP-forwarding gateways connected by a /30 WAN.

```
# Basic topology, not shell commands
LAN-A host      10.0.1.2/24    (ns: hA,  vAHost)
Gateway-A LAN   10.0.1.1/24    (ns: gwA, vAGwLan)
Gateway-A WAN   192.0.2.1/30   (ns: gwA, vAGwWan)
Gateway-B WAN   192.0.2.2/30   (ns: gwB, vBGwWan)
Gateway-B LAN   10.0.2.1/24    (ns: gwB, vBGwLan)
LAN-B host      10.0.2.2/24    (ns: hB,  vBHost)
```

1.1 Create the namespaces

Namespaces isolate interfaces, routes, and firewall rules so each element acts like an independent machine.

```
# Acquire sudo privileges
sudo -s
# Create namespaces
for ns in hA gwA gwB hB; do sudo ip netns add $ns; done
# Verify
ip netns list    # should print hA gwA gwB hB
```

1.2 Create three veth cables

Each veth pair is a zero-latency Ethernet cable inside the RAM. Three are needed: LAN-A ↔ GW-A, GW-A ↔ GW-B, LAN-B ↔ GW-B.

```
# Create the veth pairs
ip link add vAHost type veth peer name vAGwLan      # LAN-A <-> GW-A
ip link add vAGwWan type veth peer name vBGwWan      # GW-A <-> GW-B (WAN)
ip link add vBGwLan type veth peer name vBHost       # GW-B <-> LAN-B
# Verify the pairs
ip link show type veth | awk '{print $1}'    # should print all pairs in both ways
```

1.3 Move veth ends into their namespaces

Interfaces start life in the default (root) namespace; it is necessary to attach each end to its host.

```
# Attach the ends of the veth cables
ip link set vAHost netns hA
ip link set vAGwLan netns gwA
ip link set vAGwWan netns gwA
ip link set vBGwWan netns gwB
ip link set vBGwLan netns gwB
ip link set vBHost netns hB
# Visibility checks should all print interface details
ip -n hA link show vAHost
ip -n gwA link show vAGwLan
ip -n gwA link show vAGwWan
ip -n gwB link show vBGwWan
ip -n gwB link show vBGwLan
ip -n hB link show vBHost
```

1.4 Bring up interfaces and assign IPs

Linux leaves new links DOWN, and routing fails if that is left as is.

```
# Bring up loopbacks
for ns in hA gwA gwB hB; do ip -n $ns link set lo up; done
```

IP addresses must be assigned to all the devices according to the basic network topology.

```
# LAN-A
ip -n hA addr add 10.0.1.2/24 dev vAHost
ip -n hA link set vAHost up
ip -n hA route add default via 10.0.1.1
# Gateway-A
ip -n gwA addr add 10.0.1.1/24 dev vAGwLan
ip -n gwA addr add 192.0.2.1/30 dev vAGwWan
ip -n gwA link set vAGwLan up
ip -n gwA link set vAGwWan up
# Gateway-B
ip -n gwB addr add 192.0.2.2/30 dev vBGwWan
ip -n gwB addr add 10.0.2.1/24 dev vBGwLan
ip -n gwB link set vBGwWan up
ip -n gwB link set vBGwLan up
# LAN-B
ip -n hB addr add 10.0.2.2/24 dev vBHost
ip -n hB link set vBHost up
ip -n hB route add default via 10.0.2.1
```

The addresses and state for the interfaces of all namespaces have been added and configured. They should all be in *UP* state.

```
ip -n hA -br addr show # should print its one UP interface
ip -n gwA -br addr show # should print its two UP interfaces
ip -n gwB -br addr show # should print its two UP interfaces
ip -n hB -br addr show # should print its one UP interface
```

1.5 Static routes across the WAN

Each gateway must be instructed on how to reach the opposite LAN.

```
# Assign routes
ip -n gwA route add 10.0.2.0/24 via 192.0.2.2 dev vAGwWan
ip -n gwB route add 10.0.1.0/24 via 192.0.2.1 dev vBGwWan
# Show routing tables
ip -n gwA route # should print the added route
ip -n gwB route # should print the added route
```

1.6 Add WAN delay, jitter, and loss

Simulates an 80-ms RTT Internet link with light loss to make the WAN connection similar to one in a real infrastructure.

```
# Add non-idealities
ip netns exec gwA tc qdisc add dev vAGwWan root netem delay 40ms 5ms loss 0.5%
ip netns exec gwB tc qdisc add dev vBGwWan root netem delay 40ms 5ms loss 0.5%
# Verify the addition of non-idealities
ip netns exec gwA tc -s qdisc show dev vAGwWan
ip netns exec gwB tc -s qdisc show dev vBGwWan
```

1.7 Plain-text connectivity test

Confirm routing and non-idealities settings.

```
# Simple ping operation to assess the connectivity
ip netns exec hA ping -c3 10.0.2.2
ip netns exec hB ping -c3 10.0.1.2
# Stress test to verify the non-idealities of the WAN connection
ip netns exec hB iperf3 -s &
# On another shell
ip netns exec hostA iperf3 -c 10.0.2.2 -u -b 20M -t 10
```

3 strongSwan (IKEv2/IPsec)

3.1 Secure channel configurations

The StrongSwan configuration files must be created for both *Host A* and *Host B*. The high-level idea is to define one **IKE_SA** (i.e., s2s) and one **CHILD_SA** (i.e., net-net) on each side. The subnets of the namespaces are mapped onto encryption policies that must be honored at connection time.

Create the following configuration files at the indicated paths (create eventually missing folders). The separation of the files into two distinct paths mirrors how the configurations would live on the two separate hosts. Each gateway specifies its local and remote interfaces across the WAN connection. The child SAs refer to the two LANs, as demonstrated by the addresses.

Note that these two files are a possible example, but the configuration can be customized according to what is intended to be tested.

Gateway A /etc/swanctl-gwA/swanctl.conf

```
connections {
    s2s {
        version = 2
        proposals = aes256-sha256-modp2048
        local_addrs = 192.0.2.1
        remote_addrs = 192.0.2.2
        local {
            auth = psk
            id = 192.0.2.1
        }
        remote {
            auth = psk
            id = 192.0.2.2
        }
        children {
            net-net {
                local_ts = 10.0.1.0/24
                remote_ts = 10.0.2.0/24
                esp_proposals = aes256gcm16-sha256
            }
        }
    }
}
secrets {
    ike-psk {
        id-1 = 192.0.2.1
        id-2 = 192.0.2.2
        secret = "secret"
    }
}
```

Gateway B /etc/swanctl-gwB/swanctl.conf

```
connections {
    s2s {
        version = 2
        proposals = aes256-sha256-modp2048
        local_addrs = 192.0.2.2
        remote_addrs = 192.0.2.1
        local {
            auth = psk
            id = 192.0.2.2
        }
        remote {
            auth = psk
            id = 192.0.2.1
        }
        children {
            net-net {
```

```

        local_ts = 10.0.2.0/24
        remote_ts = 10.0.1.0/24
        esp_proposals = aes256gcm16-sha256
    }
}
secrets {
    ike-psk {
        id-1 = 192.0.2.1
        id-2 = 192.0.2.2
        secret = "secret"
    }
}

```

3.2 StrongSwan configuration files

Instantiate two **StrongSwan configuration files** to instruct the *charon* daemon to look for a different base configuration compared to the default one (which would be at `/etc/strongswan.conf`). The default VICI socket would be at `/var/run/charon.vici`, but here it is replaced with a custom one for each namespace. This way, the daemons do not fight for the common one, guaranteeing isolation.

Gateway A - `/etc/netns/gwA/strongswan.conf`

```

charon {
    plugins {
        vici {
            socket = unix:///run/charon-gwA.vici
        }
    }
}

```

Gateway B - `/etc/netns/gwB/strongswan.conf`

```

charon {
    plugins {
        vici {
            socket = unix:///run/charon-gwB.vici
        }
    }
}

```

3.3 Start charon & load configurations

Launch the daemon process inside the two namespaces and create a namespace-local UNIX socket for VICI control (using the file created in the previous step. This way, *swanctl* is able to talk to the correct daemon without concurrency fights. Then, load the configuration file into the running daemon using the newly created VICI socket and parsing the custom directory for the file.

```

# Launch the daemons for in both gateways
ip netns exec gwA charon-systemd --nofork &
ip netns exec gwB charon-systemd --nofork &
# Load the secure channel configurations
ip netns exec gwA env SWANCTL_DIR=/etc/swanctl-gwA \
    swanctl --load-all --uri unix:///run/charon-gwA.vici
ip netns exec gwB env SWANCTL_DIR=/etc/swanctl-gwB \
    swanctl --load-all --uri unix:///run/charon-gwB.vici

```

3.4 Bring tunnel up

Send a VICI request to **initiate** the daemon and the IKE.SA. The *net-net* CHILD_SA is set up within the *s2s* profile. This is done in the *gwA* namespace, *gwB* will respond, and the connection is established,

installing two Security Associations and two Security Policies (one per direction). This operation could be performed by initiating the connection in the *gwB* namespace without differences.

```
# Initiate from gwA
ip netns exec gwA env SWANCTL_DIR=/etc/swanctl-gwA \
    swanctl --initiate --child net-net --uri unix:///run/charon-gwA.vici
```

Verify that the configurations and the SAs have been correctly installed:

```
# Print the SAs, which should recall the configuration details
sudo ip netns exec gwA swanctl --list-sas --pretty \
    --uri unix:///run/charon-gwA.vici
sudo ip netns exec gwB swanctl --list-sas --pretty \
    --uri unix:///run/charon-gwB.vici
# Print the states, which should be mirrored according to the gateway
sudo ip -n gwA xfrm state | head
sudo ip -n gwB xfrm state | head
```

3.5 Functional test

This test with **ping** can prove that only *ESP* traffic is sniffed by *tcpdump*, proving that traffic has been secured.

```
# on a shell
ip netns exec hA ping -c3 10.0.2.2
# on another shell
ip netns exec gwA tcpdump -n -i vAGwWan esp -c 6    # should print 6 ESP captures
```

3.6 Teardown

To **close the connection**, the *charon* daemons must be killed. Then, the IKE SAs must be erased from the RAM. Finally, the kernel XFRM tables must be flushed.

```
# kill the charon daemons
sudo pkill -x charon-systemd
# flush residual states and policies
for ns in gwA gwB; do
    sudo ip netns exec $ns ip xfrm state flush
    sudo ip netns exec $ns ip xfrm policy flush
done
```